

# HDF5 Troubleshooting

Gerd Heber, The HDF Group

# A Simple Problem

Writing multiple 2D array variables over time:

ACROSS  $P$  processes arranged in a  $R \times C$  process grid

```
FOREACH step 1 .. S
```

```
    FOREACH count 1 .. A
```

```
        CREATE a double ARRAY of size  $[X,Y]$  |  $[R*X,C*Y]$  (strong | weak)
```

```
        (WRITE | READ) the ARRAY (to | from) an HDF5 file
```

```
    END
```

```
END
```

```
END
```

$S(\text{steps}) = 20$ ,  $A(\text{rrays}) = 500$ ,  $X = 100$ ,  $Y = 200$  (See [adios\\_iotest](#))




Figure: [GeeksForGeeks](#)



# Missing Information

- How are the array variables represented in HDF5?
  - 2D, 3D, 4D datasets
  - Are the extents known a priori?
  - How are the dimensions ordered?
  - Groups?
- How (order) is the data written and is the data read the same way?
- What's that storage layout?
  - How many physical files?
  - Contiguous or chunked, etc.
  - Is the data compressible?
- What's the file system or data store?
- Collective vs. independent MPI-IO
- ...

# Basic Combinations (24)

- Six griddings

- `/step=[0..10]/array=[0..499]`
- `/array=[0..499]/step=[0..19]`
- `/step=[0..20]`
- `/array=[0..499]`
- `/dataset`
- `/dataset`

Dataset {100, 200}

Dataset {100, 200}

Dataset {500, 100, 200}

Dataset {20, 100, 200}

Dataset {20, 500, 100, 200}

Dataset {500, 20, 100, 200}

- Two layouts

- Contiguous or chunked

- Two MPI modes

- Collective or independent

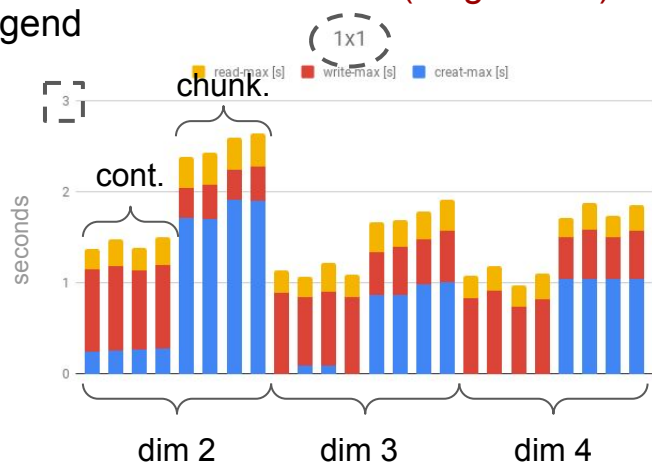
- ...

# Baseline

- Run **24 parameter configurations**
- Weak scaling
  - Each process writes  $500 \times 100 \times 200 \times 8$  (~ 80 MB) per step (20 steps)
- Single processor, 4 processor grids: 1 x 4, 2 x 2, 4 x 1
- Measure times for dataset creation, write, and read

Taller bar "=" bad (longer time)


Legend



Shorter bar "=" good (shorter time)


4x1

read-max [s] write-max [s] creat-max [s]




1x4

read-max [s] write-max [s] creat-max [s]



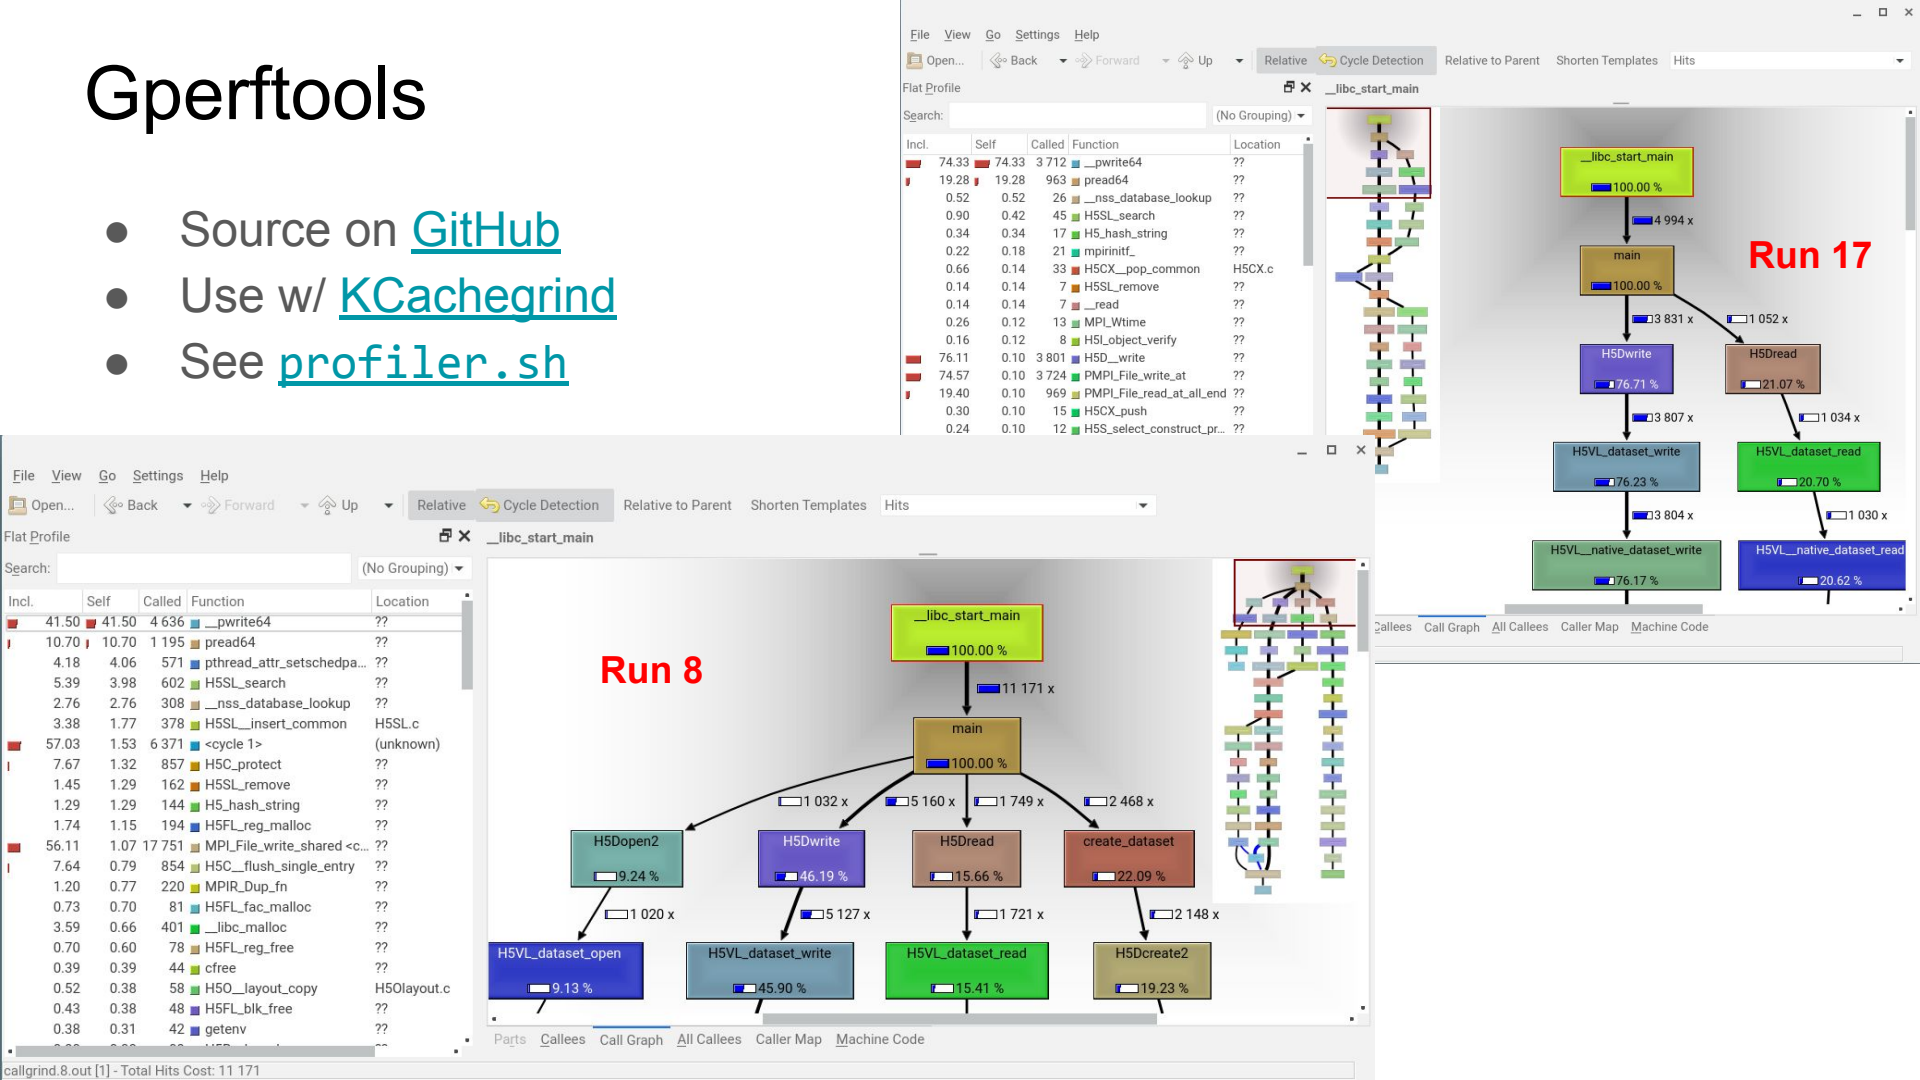
2x2

read-max [s] write-max [s] creat-max [s]



# Gperftools

- Source on [GitHub](#)
- Use w/ [KCachegrind](#)
- See [profiler.sh](#)



## 4x1 Run 7

