

Hands-on Parallel HDF5 Tutorial

October 13, 2020



© The HDF Group.

Scot Breitenfeld and Elena Pourmal, The HDF Group
Quincey Koziol, NERSC

Tutorial Prerequisites

- Tutorial assumes basic knowledge of
 - HDF5
 - MPI and MPI I/O
 - Concept of independent and collective I/O
 - C
- Access to HPC resource is provided by The HDF Group

Tutorial Goals

- Give a quick introduction to parallel HDF5 capabilities
- Help to avoid common mistakes when using parallel HDF5 library
 - For more information, check The HDF Group support portal and recent Webinars
 - <https://portal.hdfgroup.org/display/HDF5/HDF5>
 - <https://portal.hdfgroup.org/display/HDF5/Introduction+to+Parallel+HDF5>
 - <https://www.hdfgroup.org/category/webinar/>
- Provide an opportunity to explore state-of-the-art HPC cluster

Tutorial Outcome

- You should have a better understanding of
 - How to use parallel HDF5 library
 - When to use collective or independent raw data I/O
 - How to avoid HDF5 “metadata storm”

Tutorial programs and how to run them

- Get Tutorial examples
git clone <https://github.com/HDFGroup/Tutorial.git>
- README.txt has info about each example program
- To compile examples
`spack load --first hdf5`
`make`
- To run examples
`srun -n 4 <exec_name>`

Let's start!

Parallel HDF5 library capabilities

- Raw data I/O
 - All MPI ranks in communicator can write and read to / from the same or different datasets
- HDF5 metadata operations
 - All ranks in MPI communicator can access all objects in the HDF5 file
- Parallel HDF5 library limitations
 - Can't create or write variable-length data
 - Doesn't support SWMR mode
 - Doesn't support independent modifications of HDF5 file structure (HDF5 metadata)

It is easy to start using parallel HDF5: h5_ex0.c and h5par_ex0.c




- Steps to convert HDF5 application to use the parallel HDF5 library
 - Convert an application to use MPI
 - Update H5Fcreate/open to use HDF5 *MPI I/O* virtual file driver (VFD)
 - Use h5pcc compilation script to build application
 - (found in the bin directory of parallel HDF5 library installation)
 - Run as you would run parallel applications on your HPC system
- Try it:
 - Compare the content of h5_ex0.c and h5par_ex0.c with `diff` command
 - Run both examples
 - Use h5dump utility to examine SDS.h5 and SDSpar.h5
- Points to remember
 - There is no difference in HDF5 files created by the sequential or parallel HDF5 library
 - Use HDF5 MPI I/O VFD to write / read HDF5 file in parallel

Writing HDF5 dataset by columns: h5_ex1.c, h5par_ex1a.c and h5par_ex1b.c




- How to convert HDF5 application to use parallel HDF5 library
 - Assign data to each MPI rank (vs. looping through parts of data to write / read) using HDF5 hyperslab selection
- Try it:
 - Compare the content of h5_ex1.c and h5par_ex1a.c
 - Notice that each hyperslab represents noncontiguous selection in the file
 - Run both examples
 - Compare the content of h5_ex1a.c and h5par_ex1b.c
 - Run h5par_ex1b.c and compare the results
- Points to remember
 - Default properties may not be optimal when doing raw data I/O
 - Consider using `H5Pset_dxpl_mpio` and `H5FD_MPIO_COLLECTIVE` to set collective I/O for `H5Dwrite/read`



⋮

230,000 rows


⋮




Writing HDF5 dataset by columns using chunking and compression

`h5par_ex1c.c` and `h5par_ex1d.c`

- Try it:
 - Compare the content of `h5_ex1a.c` and `h5par_ex1c.c`
 - Run `h5par_ex1c.c`
 - Compare the content of `h5_ex1c.c` and `h5par_ex1d.c`
 - See how to set compression for parallel applications
 - Run `h5par_ex1d` and check the sizes of `h5par_ex1c.h5` and `h5par_ex1d.h5`
- Points to remember
 - I/O is performed on the *whole* chunk
 - Use chunks of at least 10-100MB (limit for chunk size is 4GB)
 - Think about chunk sizes if data must be read back
 - Compression requires *collective* raw data I/O



Creating HDF5 file structure: h5par_ex2a.c and h5par_ex2b.c



- Try it:
 - Let's try to create the file structure shown here
 - Compare the content of h5_ex2a.c and h5par_ex2b.c
 - Run h5par_ex2a and h5par_ex2b
 - Run h5dump -H h5par_ex2a.h5 and h5dump -H h5par_ex2b.h5
- Points to remember
 - Any operations that change HDF5 file structure or HDF5 metadata requires collective operations
 - See HDF5 documentation, General Topics in HDF5, Parallel HDF5
<https://portal.hdfgroup.org/display/HDF5/Collective+Calling+Requirements+in+Parallel+HDF5+Applications>

Reading HDF5 file structure: h5par_ex2c.c



- Try it:
 - See h5par_ex2c.c for how to read HDF5 metadata independently
 - Run h5par_ex2c
- Points to remember
 - H5Fopen is *always a collective call*
 - H5G/D/Aopen calls *can be independent* if the file structure is not modified

Using independent H5Dopen to write data from each MPI rank:

`h5par_ex2d.c` `h5par_ex2e.c`

- Try it:
 - Review `h5par_ex2d.c` and `h5par_ex2e.c` for how to create datasets and to write them independently from each MPI rank
 - Run `h5par_ex2d` and `h5par_ex2e`
 - Use `h5dump h5par_ex2d.h5` to check the content
- Points to remember
 - Use `H5Pset_alloc_time` dataset creation property and `H5D_ALLOC_TIME_EARLY` to allocate space in the file
 - Shown approach cannot be applied to *extensible* and *compressed* datasets

Avoid MD storm writes and reads: h5par_comparison

- Try it:
 - Review `h5par_comparison.c`; it creates a dataset and writes it with hyperslabs using all ranks
 - Notice that multiple ranks access the **same** HDF5 metadata when creating a file and a dataset, and again when opening the file and the dataset. (Performance may suffer!)
 - Run `h5par_comparison`, `h5par-comparison-collmd` and `h5par-comparison-collio`
- Points to remember
 - When all ranks access the same objects in HDF5 file avoid “metadata write/read storm” by using `H5Pset_all_coll_metadata_ops` and `H5Pset_coll_metadata_write`
 - See <https://portal.hdfgroup.org/display/HDF5/General+Access+Properties>
 - Other hints:
 - Use HDF5 1.10.7 and later to use optimizations when reading the entire dataset by all MPI ranks
 - Use compact storage for a small dataset (order of KBs) and when reading/writing the dataset by all MPI ranks along with optimization above to avoid metadata storm

THANK YOU!

Questions & Comments?